



COMSATS University Islamabad (Lahore Campus)

Assignment <2> FALL 2022

Course Title:	Computational Intelligence	Course Code:	CSC	Credit Hours:	3(2,1)
Course Instructor:	Dr. Atifa Athar	Programme Name:	BS Computer Science		
Semester:	7 th	Batch:	FA19	Section:	B
Due Date:	04-11-2022	Maximum Marks:	10		
Student Name:	Muhammad Awais Abdullah Muhammad Huzaifa Tariq	Registration No.	FA19-BCS-105 FA19-BCS-114 FA19-BCS-047		

Important Instructions / Guidelines:

- No late submissions will be accepted.
- All assignments are required to be submitted using attached template only.

Question No 1.

Marks: 10

CLO: <3>; Bloom Taxonomy Level: <Applying>

Select a textual dataset of your choice (for example IMDB movies review data). Design and implement RNN and LSTM Neural Network, Bert model (along with its variants) in Python using your selected dataset.

Submit a complete report including an introduction to the dataset. implementation, Results (accuracy tables and confusion matrix), and comparisons. Along with the Report, you are required to submit your code as well.

Problem Statement:

Design, implement, evaluate and compare different NLP techniques for tweet classification such as RNN & LSTM, CNN, CNN+LSTM, and BERT in python on the COVID-19 vaccines dataset.

Introduction to Dataset:

The dataset we have selected for this assignment is "Tweets Dataset on COVID-19 Vaccines". It was collected by social media posts on COVID-19 vaccines using Twitter's streaming API from December 9, 2020 - Feb 24, 2021, with keywords related to the vaccines ("Vaccine", "Pfizer", "BioNTech", "Moderna", "Janssen", "AstraZeneca", "Sinopharm"). This dataset includes CSV files that contain IDs, text, labels, and sixty other columns related to the COVID-19 tweets. On social media, content propagates through the network when accounts engage with posts by re-sharing (retweeting), replying to tweets, and quoting (quote tweets are retweets without a comment). There are 376K entries in the raw dataset with most of the tweets as retweeted, so there are unique 14.6k tweets, roughly 10,377 labeled as reliable and 4,267 labeled as unreliable/conspiracy.

Data Info:

RangeIndex: 14644 entries, 0 to 14643

Data columns (total 69 columns):

#	Column	Non-Null Count	Dtype
---	-----	-----	-----
0	cascade_sno	14644 non-null	int64
1	cascade_tid	14644 non-null	int64
2	tweetid	14644 non-null	int64
3	userid	14644 non-null	int64
4	screen_name	14644 non-null	object
5	date	14644 non-null	object
6	lang	14644 non-null	object
7	location	10398 non-null	object
8	place_id	52 non-null	object
9	place_url	52 non-null	object
10	place_type	52 non-null	object
11	place_name	52 non-null	object
12	place_full_name	52 non-null	object
13	place_country_code	52 non-null	object
14	place_country	52 non-null	object
15	place_bounding_box	52 non-null	object
16	text	14644 non-null	object
17	extended	14644 non-null	object
18	coord	0 non-null	float64
19	reply_userid	388 non-null	float64
20	reply_screen	388 non-null	object
21	reply_statusid	388 non-null	float64
22	tweet_type	14644 non-null	object
23	friends_count	14644 non-null	int64
24	listed_count	14644 non-null	int64
25	followers_count	14644 non-null	int64
26	favourites_count	14644 non-null	int64
27	statuses_count	14644 non-null	int64
28	verified	14644 non-null	bool
29	hashtag	14644 non-null	object
30	urls_list	14644 non-null	object
31	profile_pic_url	12200 non-null	object
32	profile_banner_url	12200 non-null	object
33	display_name	14643 non-null	object
34	date_first_tweet	14644 non-null	object
35	account_creation_date	14644 non-null	object
36	rt_urls_list	14644 non-null	object
37	mentionid	14644 non-null	object
38	mentionsn	14644 non-null	object
39	rt_screen	8821 non-null	object
40	rt_userid	8821 non-null	float64
41	rt_text	8889 non-null	object
42	rt_hashtag	14644 non-null	object
43	rt_qtd_count	14644 non-null	int64
44	rt_rt_count	14644 non-null	int64
45	rt_reply_count	14644 non-null	int64
46	rt_fav_count	14644 non-null	int64
47	rt_tweetid	8889 non-null	float64
48	rt_location	7181 non-null	object
49	qtd_screen	129 non-null	object
50	qtd_userid	129 non-null	float64
51	qtd_text	132 non-null	object

52	qtd_hashtag	14644	non-null	object
53	qtd_qtd_count	14644	non-null	int64
54	qtd_rt_count	14644	non-null	int64
55	qtd_reply_count	14644	non-null	int64
56	qtd_fav_count	14644	non-null	int64
57	qtd_tweetid	132	non-null	float64
58	qtd_urls_list	14644	non-null	object
59	qtd_location	97	non-null	object
60	sent_vader	14644	non-null	float64
61	token	14641	non-null	object
62	media_urls	14644	non-null	object
63	rt_media_urls	14644	non-null	object
64	q_media_urls	14644	non-null	object
65	corrected_tweet_type	14644	non-null	object
66	ns_label	14644	non-null	object
67	ns_url	14644	non-null	object
68	timestamp	14644	non-null	float64

Data Preprocessing:

1. Clean all the tweets by converting them into lowercase and remove all unnecessary characters (and some unnecessarily repeated words).
2. Encode the string labels as “1” for unreliable/conspiracy and “0” for reliable.
3. Split the data into training and testing sets.
4. Generate tokens of the tweets data.
5. Generate sequences and add padding to the max length.

NLP techniques for Text Classification:

There are several natural language processing (NLP) techniques that can be used for text classification:

1. **Bag-of-Words:** This technique represents each text as a bag of its words, disregarding grammar and word order but keeping track of occurrences. This can be used as input for a classifier such as Naive Bayes.
2. **Term Frequency-Inverse Document Frequency (TF-IDF):** This technique assigns weight to each word in the text, with more weight given to words that are more informative and less common words. This can also be used as input for a classifier.
3. **Word Embeddings:** This technique represents each word in a text as a dense vector, which captures semantic and syntactic information about the word. These vectors can be used as input for a classifier such as a feed-forward neural network or a convolutional neural network.
4. **Recurrent Neural Networks (RNNs) and Long Short-Term Memory (LSTM):** These are types of neural networks that process sequential data, such as text. They can be trained in text classification tasks and can output a probability of the text belonging to each class.
5. **Transformers:** These are deep learning models that are particularly well suited for NLP tasks, such as text classification. They have been shown to be very effective in achieving state-of-the-art performance on a wide range of NLP tasks.

LSTM

LSTM stands for Long Short-Term Memory. It is a type of Recurrent Neural Network (RNN) that can capture long-term dependencies in sequential data. Unlike traditional RNNs, which tend to forget past information as the network unfolds over time, LSTMs have a built-in memory cell that can retain information for prolonged periods. This makes them well-suited for tasks such as language modeling, machine translation, and speech recognition.

Implementation: Implementing an LSTM network typically involves the following steps:

1. **Prepare the data:** The first step is to prepare the data by tokenizing and indexing the text, and then dividing it into training and testing sets.
2. **Choose the LSTM architecture:** The next step is to choose the LSTM architecture, which includes deciding on the number of layers, the number of units in each layer, and whether to use a bidirectional or unidirectional LSTM.
3. **Define the model:** The LSTM model can be defined using a deep learning library such as TensorFlow, Keras, or PyTorch. The model consists of an LSTM layer, followed by one or more fully connected layers.
4. **Compile the model:** Once the model is defined, it needs to be compiled by specifying the loss function, optimizer, and metrics.
5. **Train the model:** The model is then trained on the training data using the `fit()` function. The model will learn to associate the input sequences with the corresponding output labels.
6. **Evaluate the model:** After the model is trained, it can be evaluated on the test data using the `evaluate()` function, which will return the loss and any other metrics that were specified during compilation.

It's important to note that the specific implementation details of an LSTM will depend on the deep learning framework and the specific use case.

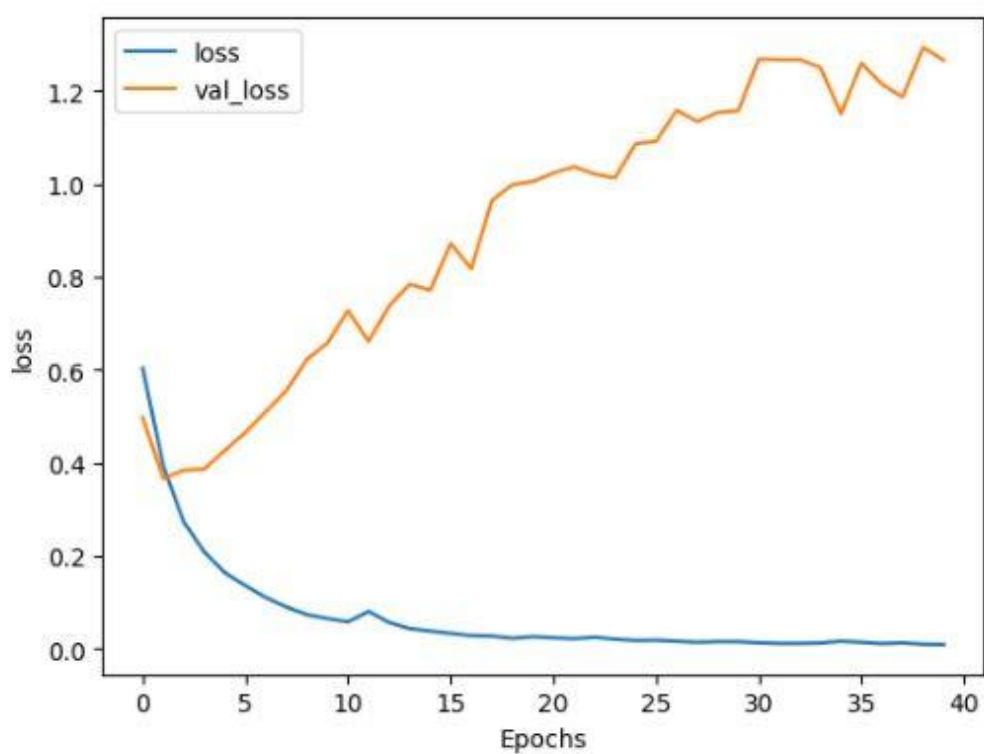
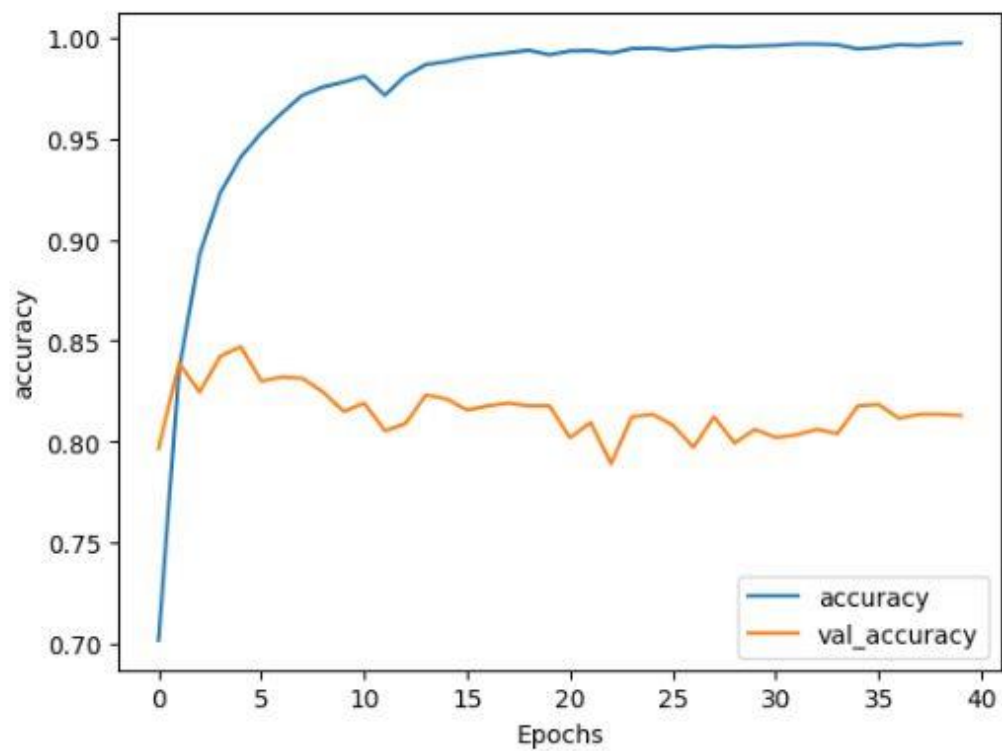
LSTM Model Summary:

Model: "sequential"

Layer (type)	Output Shape	Param #
embedding (Embedding)	(None, 331, 16)	224000
spatial_dropout1d (SpatialDropout1D)	(None, 331, 16)	0
bidirectional (Bidirectional)	(None, 64)	12544
dense (Dense)	(None, 6)	390
dense_1 (Dense)	(None, 1)	7

=====
Total params: 236,941
Trainable params: 236,941
Non-trainable params: 0
=====

	precision	recall	f1-score	support
Reliable	0.90	0.85	0.87	1119
UnReliable	0.59	0.69	0.64	346
accuracy			0.81	1465
macro avg	0.74	0.77	0.75	1465
weighted avg	0.83	0.81	0.82	1465



GRU

The Gated Recurrent Unit (GRU) is a type of recurrent neural network (RNN) architecture that addresses some of the limitations of traditional RNNs, such as the vanishing gradient problem. The GRU architecture utilizes gating mechanisms to control the flow of information through the network, allowing for better preservation of long-term dependencies and improved performance on tasks such as language modeling and speech recognition. The GRU was introduced by Cho et al. in 2014 and has since become a popular choice in RNN architectures.

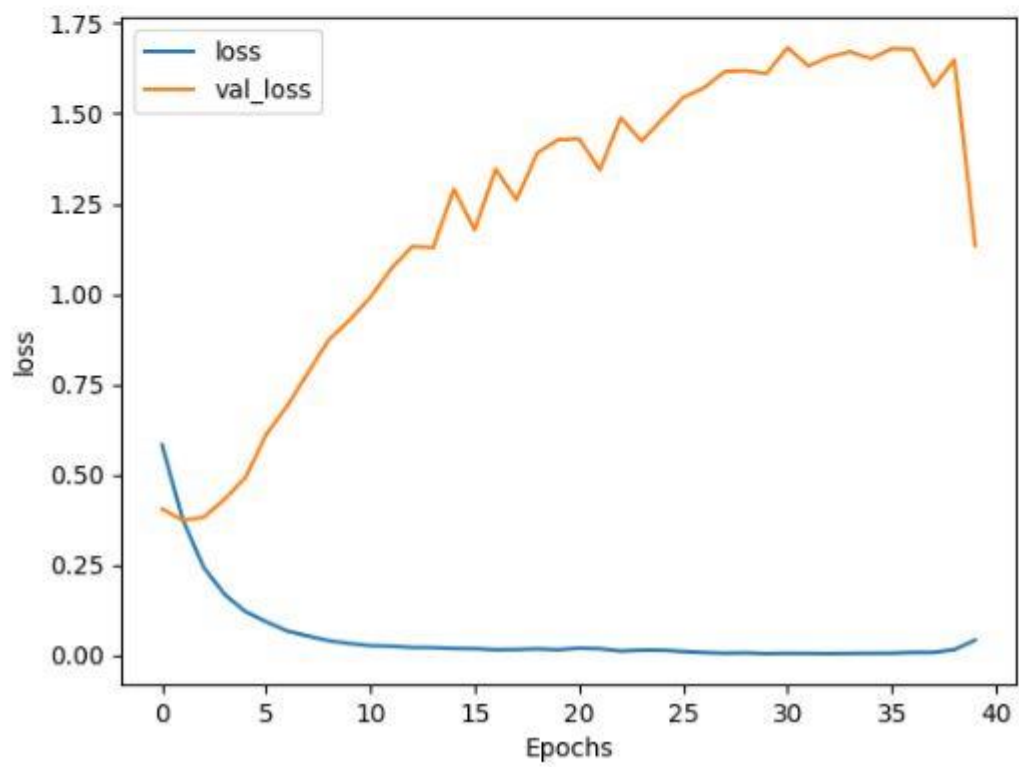
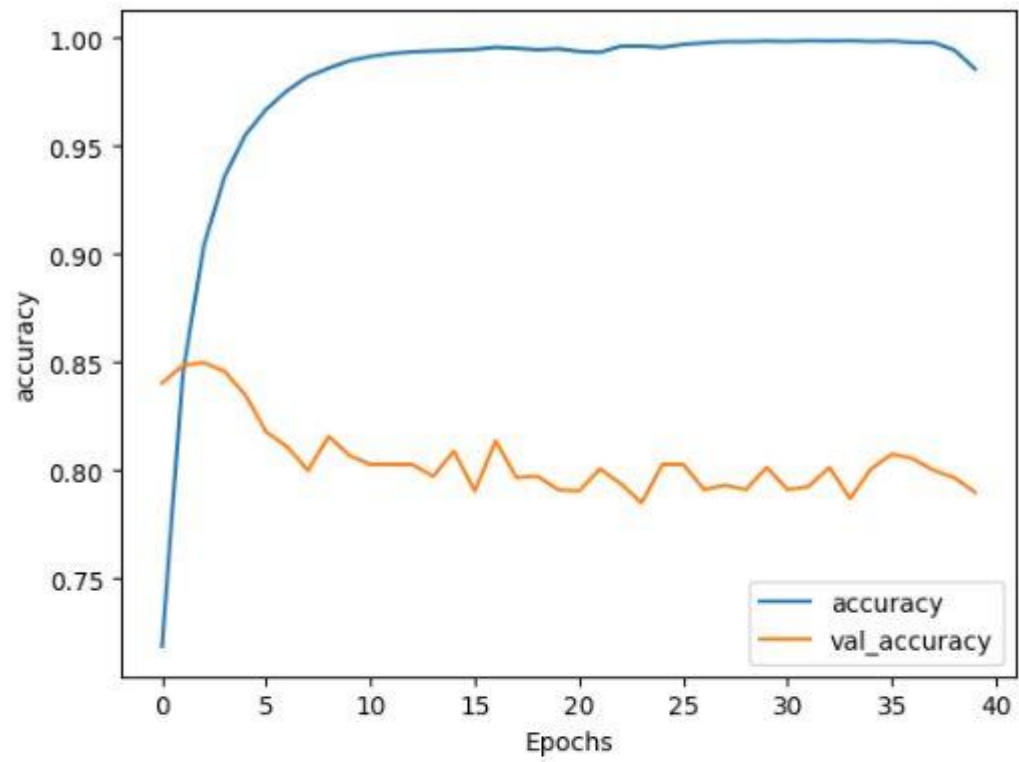
Implementation: Implementing a GRU in code involves several steps:

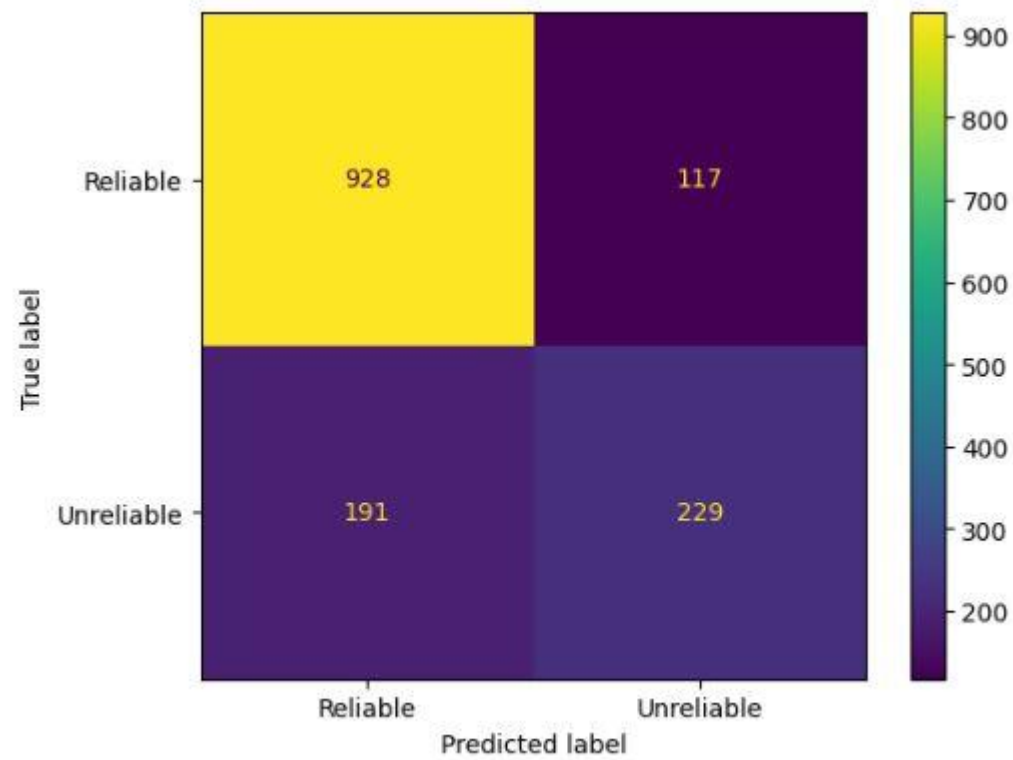
1. Initializing the GRU model. This typically includes defining the number of hidden units, the input dimension, and the number of layers in the model.
2. Defining the input and output layers. The input layer is typically an embedding layer, which converts the input words or sequences into a dense vector representation. The output layer is typically a linear layer that maps the hidden state of the GRU to the output of the model.
3. Defining the GRU layers. This typically involves creating instances of the GRU class, specifying the number of hidden units, and connecting the input and output layers to the GRU layers.
4. Training the model. This typically involves defining a loss function and an optimizer and then training the model on a dataset using a training loop.
5. Evaluating the model. This typically involves evaluating the model on a validation or test dataset and measuring its performance using metrics such as accuracy or perplexity.

GRU Model Summary:

Model: "sequential_1"

Layer (type)	Output Shape	Param #
embedding_1 (Embedding)	(None, 331, 16)	224000
bidirectional_1 (Bidirectional)	(None, 64)	9600
dense_2 (Dense)	(None, 6)	390
dense_3 (Dense)	(None, 1)	7
Total params: 233,997		
Trainable params: 233,997		
Non-trainable params: 0		





	precision	recall	f1-score	support
Reliable	0.89	0.83	0.86	1119
UnReliable	0.55	0.66	0.60	346
accuracy			0.79	1465
macro avg	0.72	0.75	0.73	1465
weighted avg	0.81	0.79	0.80	1465

CNN

A Convolutional Neural Network (CNN) is a type of deep learning neural network that is designed to process and analyze data that has a grid-like structure, such as images, videos, and audio signals. CNNs are composed of multiple layers, including convolutional layers, pooling layers, and fully connected layers. The convolutional layers are responsible for analyzing the input data and extracting features, such as edges, textures, and patterns. The pooling layers are responsible for reducing the spatial dimensions of the data while maintaining the important features. The fully connected layers are responsible for classifying the data based on the features extracted by the convolutional and pooling layers.

CNNs are particularly useful for image classification, object detection, and image segmentation tasks, as they are able to learn and extract features from images that are not easily represented by traditional image processing techniques. They have been widely adopted in many computer vision applications such as self-driving cars, image search engines, and facial recognition. CNNs can also be used for text classification tasks, in which the goal is to classify a piece of text into one or more predefined categories. Text classification is a widely used technique in natural language processing (NLP) and is used in tasks such as sentiment analysis, topic classification, and spam detection. In text classification, a CNN is typically applied to the embedding of words in a sentence or a document, rather than to the raw input text. The embedding is a dense vector representation of each word, which captures the semantic meaning of the word. The embedding is typically obtained by training a word embedding model, such as word2vec or GloVe, on a large corpus of text.

The CNN architecture for text classification is similar to the CNN architecture for image classification, but with some modifications to adapt it to text data. The convolutional layers are typically 1-dimensional (1D) convolutional layers, which are applied to the embedding of the words in a sentence or a document. The pooling layers are typically max-pooling layers, which are applied to reduce the spatial dimensions of the data and to maintain the important features. The fully connected layers are responsible for classifying the data based on the features extracted by the convolutional and pooling layers.

In summary, CNNs are a powerful tool for analyzing grid-like data and are particularly useful for image and video processing tasks but they can be used for text classification tasks by applying 1D convolutional layers to the embedding of words in a sentence or a document and using max-pooling and fully connected layers to classify the data based on the features extracted by the convolutional and pooling layers. It is important to note that this type of architecture, while it can work well, is not the most common way to perform text classification and other architectures such as LSTM, GRU, or Transformer architectures are commonly used instead.

Implementation: Here are the general steps for implementing a CNN for text classification:

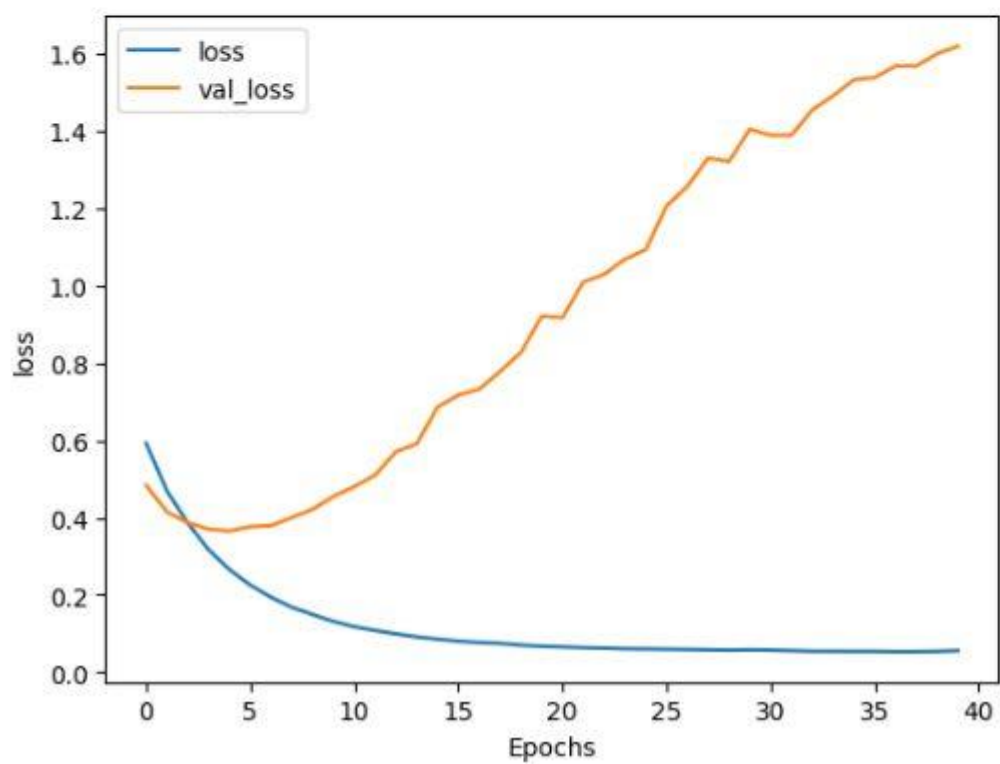
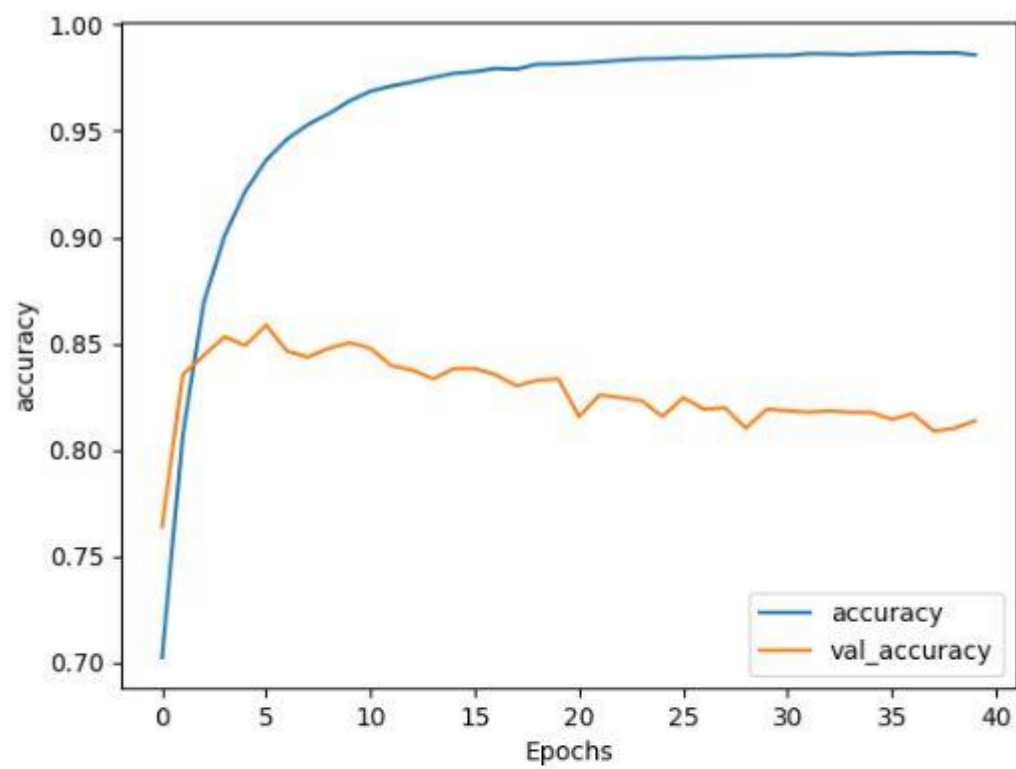
1. **Data preprocessing:** Clean and preprocess the text data, such as removing stop words, stemming, or lemmatizing the words, and tokenizing the text.
2. **Word embedding:** Train a word embedding model, such as word2vec or GloVe, on a large corpus of text, and use the trained model to obtain the embedding of the words in the text data.
3. **Define the CNN model:** Initialize the CNN model and define the architecture, including the number of convolutional layers, the number of filters, the filter size, the pooling layers, and the fully connected layers.

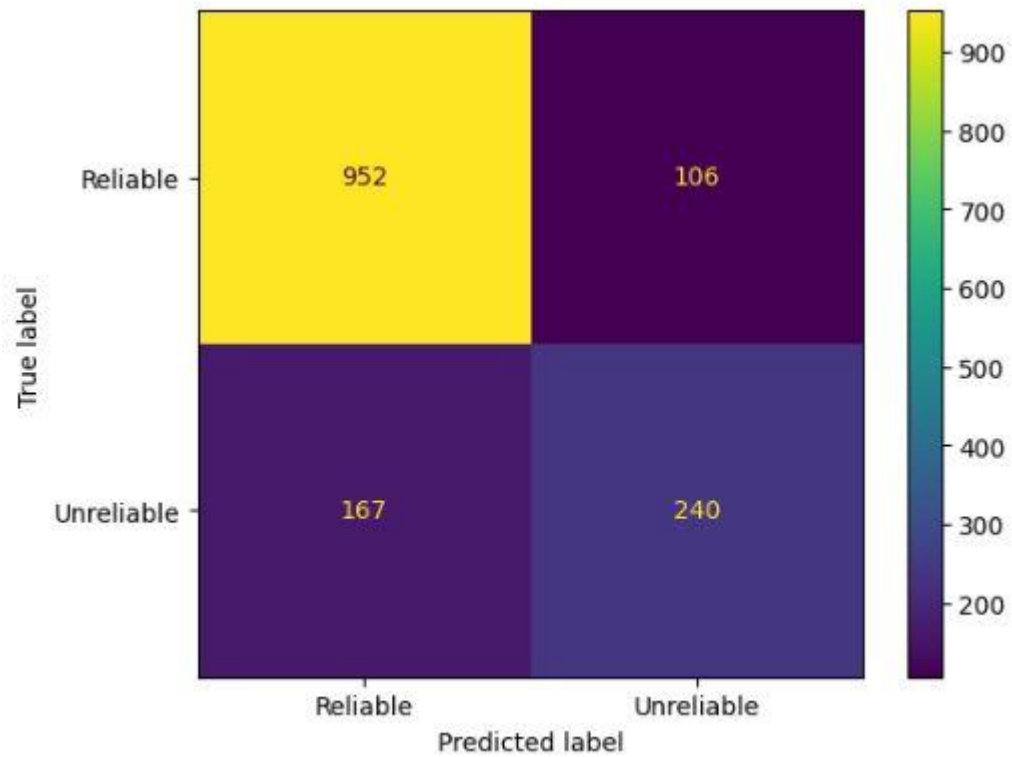
4. **Train the model:** Train the model on the text data and the corresponding labels, using a loss function such as cross-entropy loss and an optimizer such as Adam.
5. **Evaluate the model:** Evaluate the model on a validation or test dataset and measure its performance using metrics such as accuracy or F1-score.

CNN Model Summary:

Model: "sequential_2"

Layer (type)	Output Shape	Param #
embedding_2 (Embedding)	(None, 331, 16)	224000
conv1d (Conv1D)	(None, 327, 128)	10368
global_average_pooling1d (GlobalAveragePooling1D)	(None, 128)	0
dense_4 (Dense)	(None, 6)	774
dense_5 (Dense)	(None, 1)	7
Total params: 235,149		
Trainable params: 235,149		
Non-trainable params: 0		





	precision	recall	f1-score	support
Reliable	0.90	0.85	0.87	1119
UnReliable	0.59	0.69	0.64	346
accuracy			0.81	1465
macro avg	0.74	0.77	0.76	1465
weighted avg	0.83	0.81	0.82	1465

CNN + LSTM

A CNN-LSTM (Convolutional Neural Network-Long Short-Term Memory) network is a type of deep learning neural network that combines the strengths of CNNs and LSTMs. CNNs are particularly good at analyzing grid-like data, such as images and videos, and extracting features from the data, while LSTMs are particularly good at preserving long-term dependencies in sequential data, such as time series and natural language.

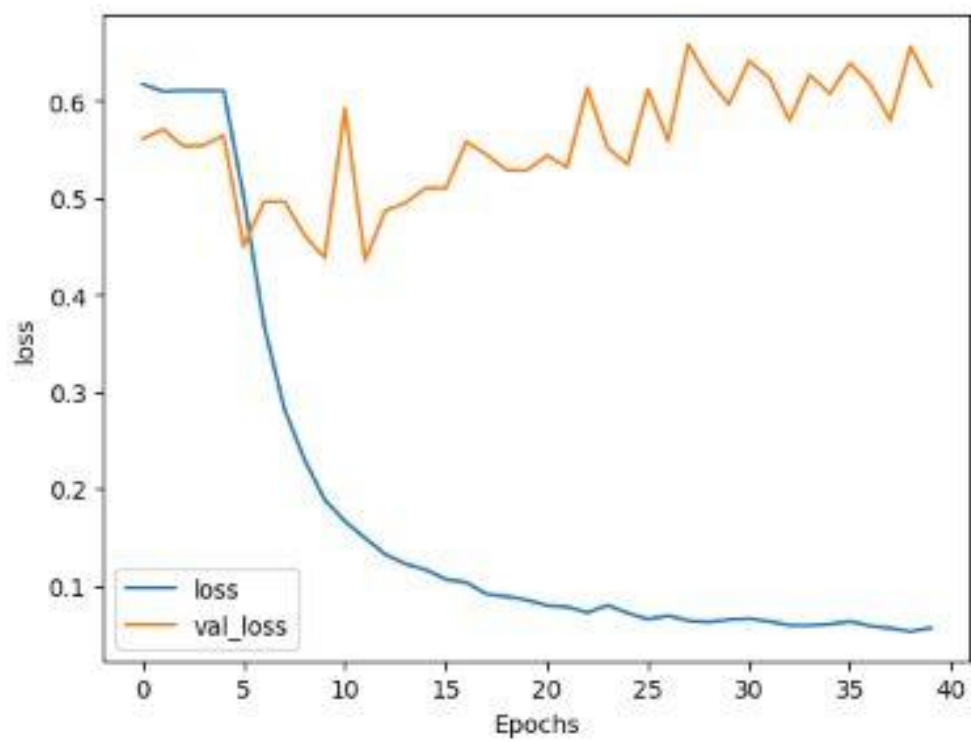
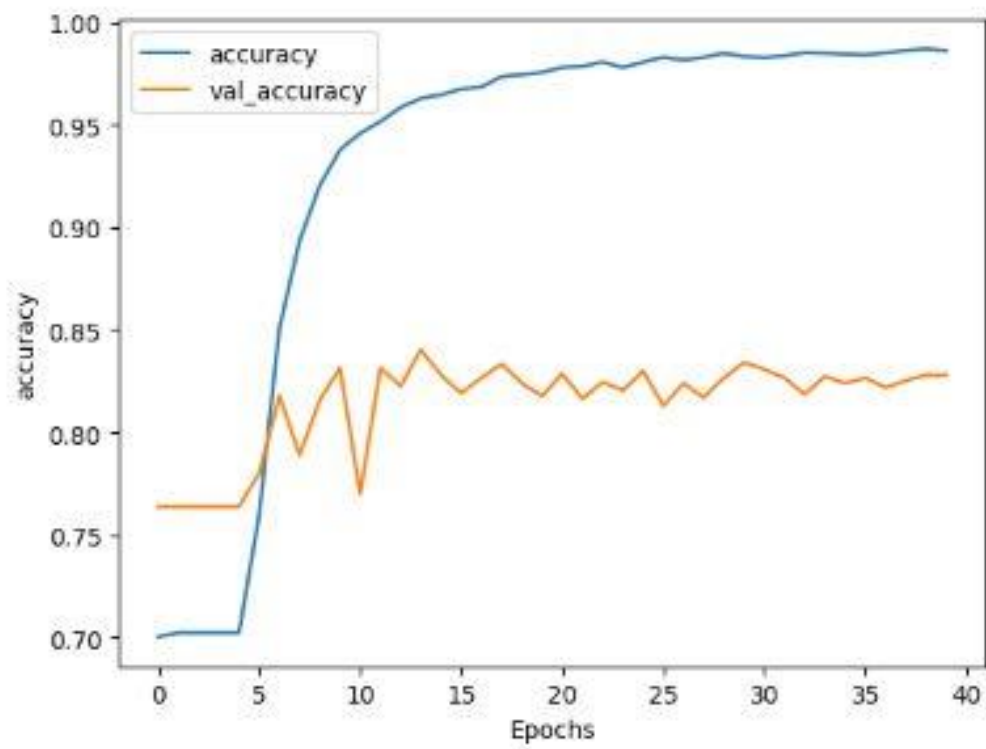
The idea behind a CNN-LSTM network is to use a CNN to extract features from the input data, and then pass the extracted features to an LSTM to preserve the long-term dependencies in the data. The CNN and LSTM are typically connected in a sequential manner, with the output of the CNN being passed as the input to the LSTM. CNN-LSTM networks are particularly useful for tasks that involve both spatial and temporal information, such as video classification and action recognition, and speech recognition.

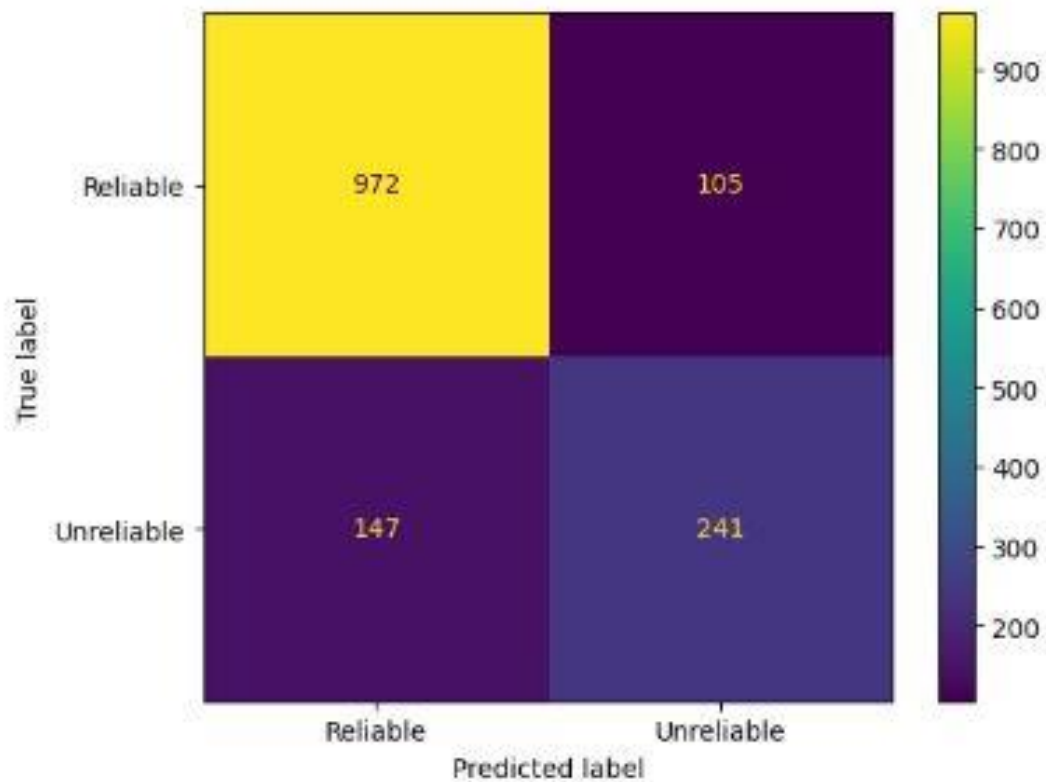
Implementation: Here are the general steps for implementing a CNN-LSTM for text classification:

1. **Data preprocessing:** Clean and preprocess the text data, such as removing stop words, stemming, or lemmatizing the words, and tokenizing the text.
2. **Word embedding:** Train a word embedding model, such as word2vec or GloVe, on a large corpus of text, and use the trained model to obtain the embedding of the words in the text data.
3. **Define the CNN-LSTM model:** Initialize the CNN-LSTM model and define the architecture, including the number of convolutional layers, the number of filters, the filter size, the pooling layers, the LSTM layers, and the fully connected layers.
4. **Train the model:** Train the model on the text data and the corresponding labels, using a loss function such as cross-entropy loss and an optimizer such as Adam.
5. **Evaluate the model:** Evaluate the model on a validation or test dataset and measure its performance using metrics such as accuracy or F1-score.

CNN Model Summary:

Model: "sequential_3"		
Layer (type)	Output Shape	Param #
embedding_3 (Embedding)	(None, 331, 16)	224000
bidirectional_2 (Bidirectional)	(None, 331, 12)	1104
conv1d_1 (Conv1D)	(None, 330, 24)	600
dropout (Dropout)	(None, 330, 24)	0
conv1d_2 (Conv1D)	(None, 328, 24)	1752
dropout_1 (Dropout)	(None, 328, 24)	0
conv1d_3 (Conv1D)	(None, 324, 24)	2904
dropout_2 (Dropout)	(None, 324, 24)	0
conv1d_4 (Conv1D)	(None, 319, 24)	3480
dropout_3 (Dropout)	(None, 319, 24)	0
conv1d_5 (Conv1D)	(None, 312, 12)	2316
dropout_4 (Dropout)	(None, 312, 12)	0
conv1d_6 (Conv1D)	(None, 309, 12)	588
dropout_5 (Dropout)	(None, 309, 12)	0
conv1d_7 (Conv1D)	(None, 306, 12)	588
dropout_6 (Dropout)	(None, 306, 12)	0
conv1d_8 (Conv1D)	(None, 304, 12)	444
dropout_7 (Dropout)	(None, 304, 12)	0
conv1d_9 (Conv1D)	(None, 302, 12)	444
dropout_8 (Dropout)	(None, 302, 12)	0
conv1d_10 (Conv1D)	(None, 301, 12)	300
dropout_9 (Dropout)	(None, 301, 12)	0
conv1d_11 (Conv1D)	(None, 300, 12)	300
dropout_10 (Dropout)	(None, 300, 12)	0
bidirectional_3 (Bidirectional)	(None, 300, 24)	2400
max_pooling1d (MaxPooling1D)	(None, 75, 24)	0
dropout_11 (Dropout)	(None, 75, 24)	0
flatten (Flatten)	(None, 1800)	0
dense_6 (Dense)	(None, 26)	46826
dense_7 (Dense)	(None, 8)	216
dense_8 (Dense)	(None, 1)	9
Total params: 288,271		
Trainable params: 288,271		
Non-trainable params: 0		





	precision	recall	f1-score	support
Reliable	0.90	0.87	0.89	1119
UnReliable	0.62	0.70	0.66	346
accuracy			0.83	1465
macro avg	0.76	0.78	0.77	1465
weighted avg	0.84	0.83	0.83	1465

BERT

BERT (Bidirectional Encoder Representations from Transformers) is a pre-trained transformer-based neural network model for natural language processing (NLP) tasks such as text classification, question answering, and named entity recognition. BERT is trained on a large corpus of text using a technique called unsupervised pre-training, which means it is trained on a massive amount of text data without any specific task in mind. This allows it to learn the general patterns and features of the language, which can then be fine-tuned for specific NLP tasks.

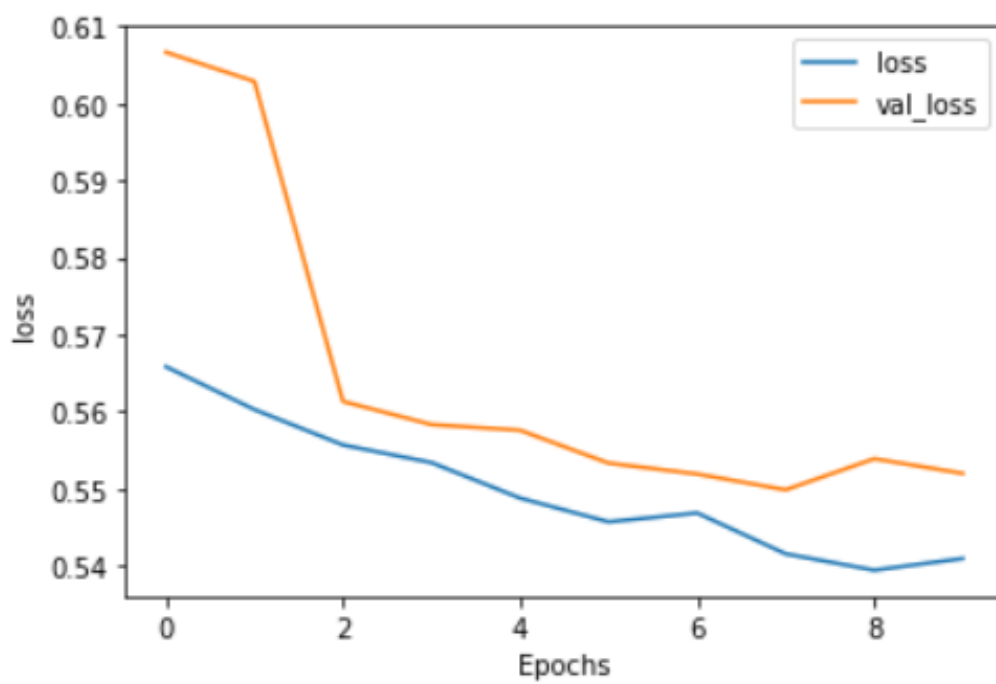
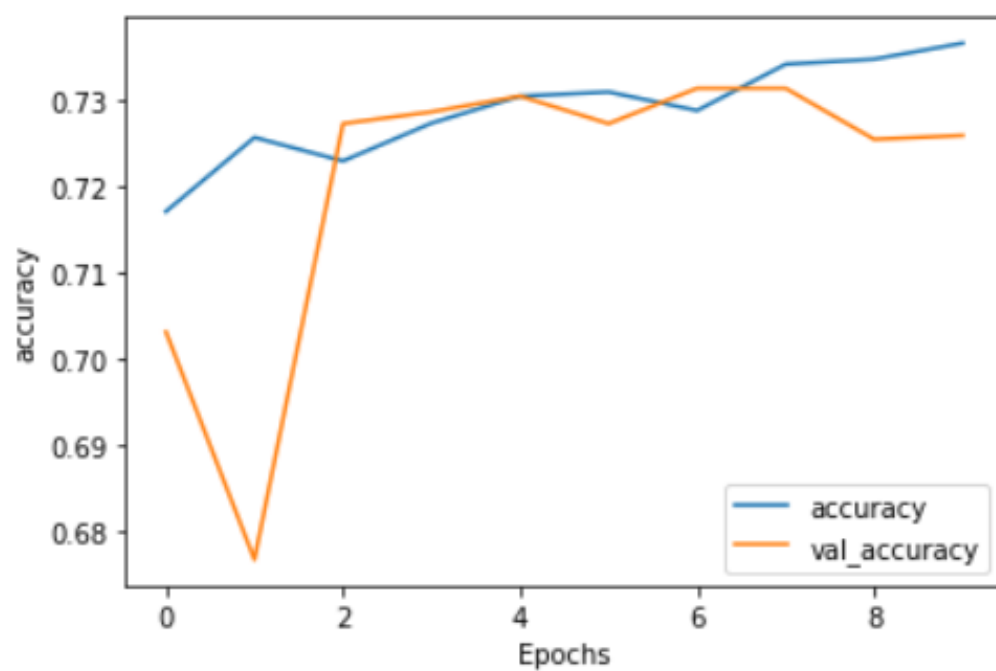
BERT is unique because it is bidirectional, meaning it considers the context of the words both to the left and to the right of the current word, rather than only the context to the left as in traditional models. This allows BERT to better understand the meaning of the words in the sentence and improve the performance on NLP tasks. BERT has achieved state-of-the-art results on a wide range of NLP tasks and has become a popular choice for natural language understanding and text classification tasks. It is available in several pre-trained versions, such as BERT-base and BERT-large, which can be fine-tuned to specific tasks and datasets.

Implementation: Here are the general steps for implementing BERT for text classification:

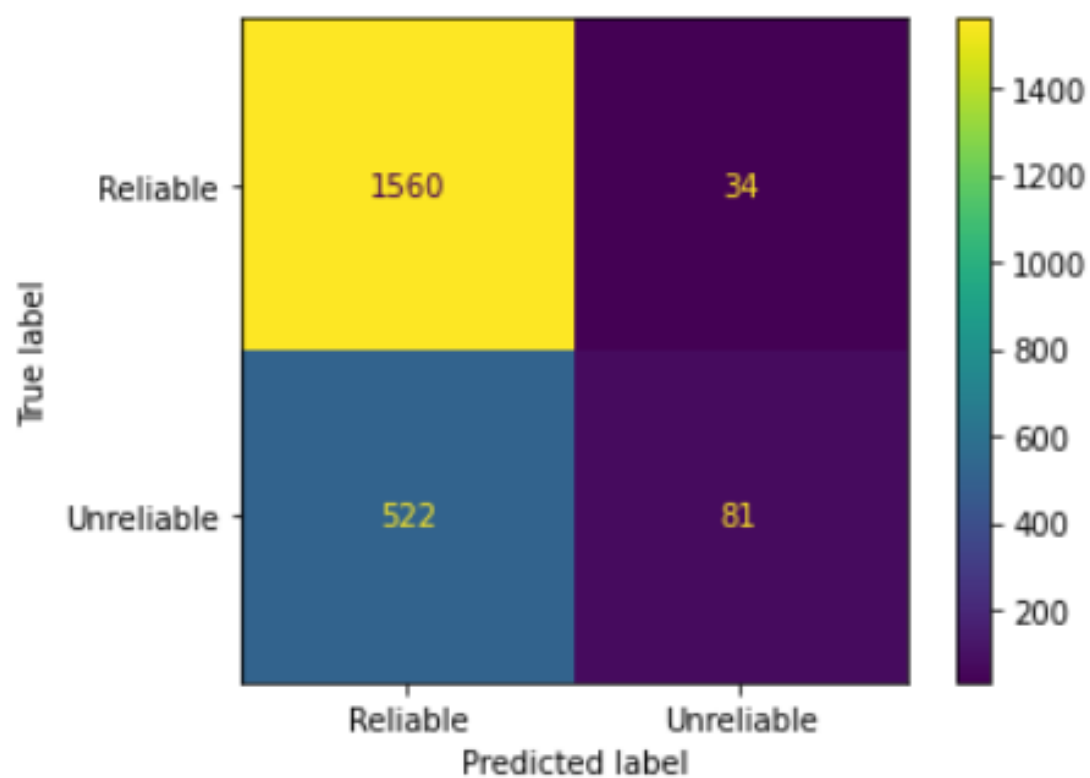
1. **Data preprocessing:** Clean and preprocess the text data, such as removing stop words, stemming or lemmatizing the words, and tokenizing the text.
2. **Download a pre-trained BERT model:** Download a pre-trained BERT model from the HuggingFace library or TensorFlow Hub, such as BERT-base or BERT-large.
3. **Tokenize the text data:** Use the BERT tokenizer to tokenize the text data, which will convert the text into a format that the BERT model can understand.
4. **Create a data loader:** Create a data loader that loads the tokenized text data and labels and creates batches for training and evaluation.
5. **Define the model architecture:** Define the architecture of the model, including the pre-trained BERT model and the additional layers for the classification task.
6. **Fine-tune the pre-trained model:** Fine-tune the pre-trained BERT model on the text classification task using a loss function such as cross-entropy loss and an optimizer such as Adam.
7. **Evaluate the model:** Evaluate the model on a validation or test dataset and measure its performance using metrics such as accuracy or F1-score.

BERT Model Summary:

Layer (type)	Output Shape	Param #	Connected to
=====			
text (InputLayer)	[(None,)]	0	[]
keras_layer (KerasLayer)	{'input_type_ids': (None, 128), 'input_word_ids': (None, 128), 'input_mask': (None, 128)}	0	['text[0][0]']
keras_layer_1 (KerasLayer)	{'sequence_output': (None, 128, 768), 'encoder_outputs': [(None, 128, 768), (None, 128, 768), (None, 128, 768), (None, 128, 768), (None, 128, 768), (None, 128, 768), (None, 128, 768), (None, 128, 768), (None, 128, 768), (None, 128, 768), (None, 128, 768), (None, 128, 768)], 'pooled_output': (None, 768), 'default': (None, 768)}	109482241	['keras_layer[0][0]', 'keras_layer[0][1]', 'keras_layer[0][2]']
output (Dense)	(None, 1)	769	['keras_layer_1[0][13]']
=====			
Total params: 109,483,010			
Trainable params: 769			
Non-trainable params: 109,482,241			
=====			



	precision	recall	f1-score	support
0	0.75	0.98	0.85	1594
1	0.70	0.13	0.23	603
accuracy			0.75	2197
macro avg	0.73	0.56	0.54	2197
weighted avg	0.74	0.75	0.68	2197



ALBERT

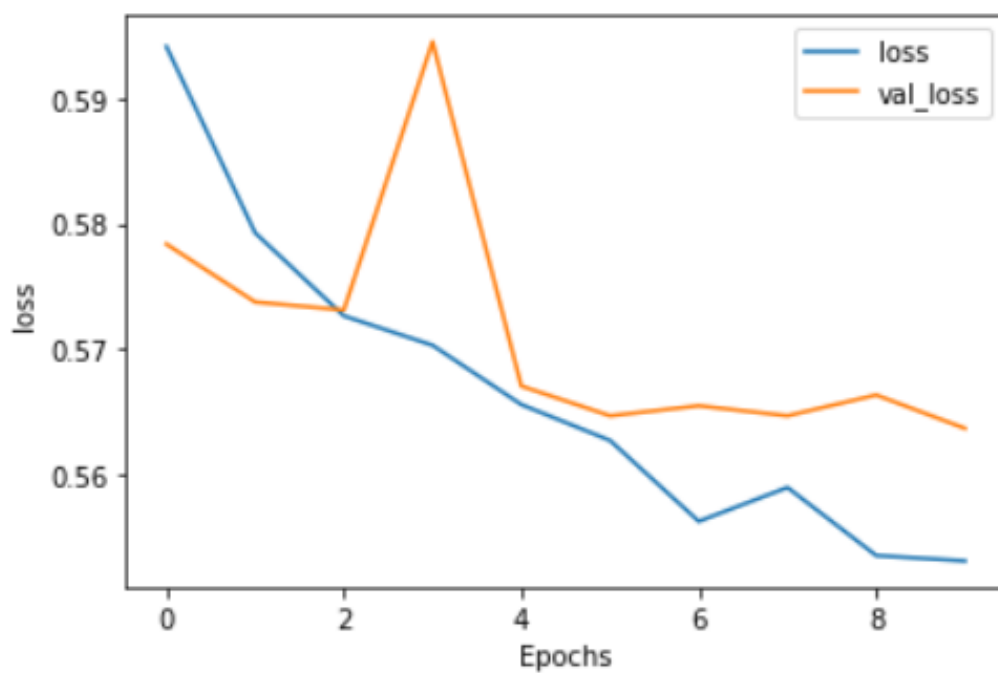
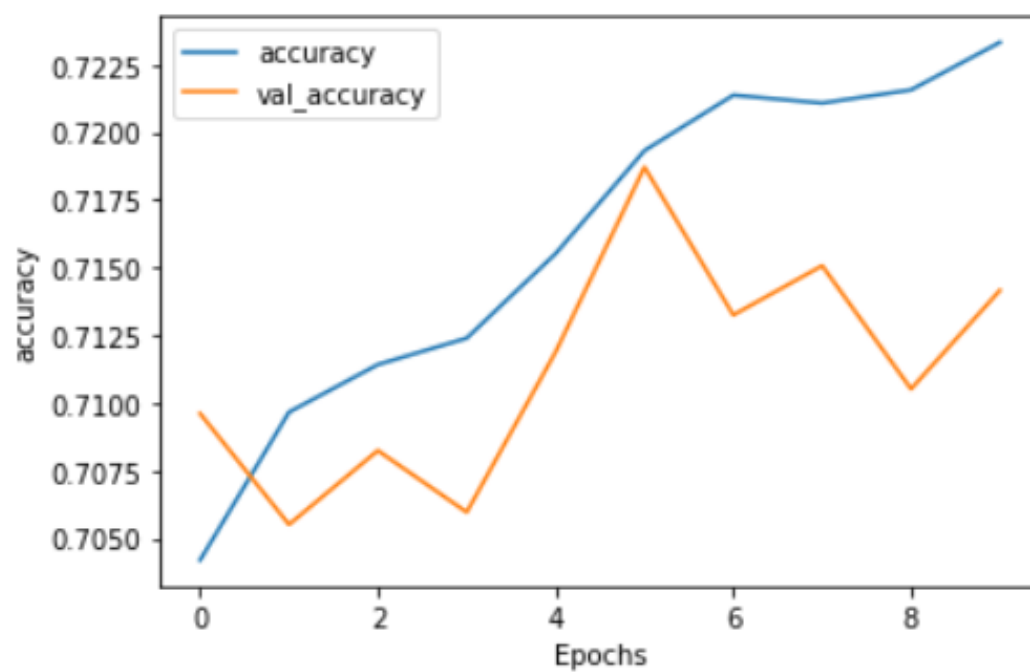
ALBERT (A Lite BERT) is a light version of BERT (Bidirectional Encoder Representations from Transformers), a pre-trained transformer-based neural network model for natural language processing (NLP) tasks such as text classification, question answering, and named entity recognition. ALBERT is similar to BERT but it is designed to be more lightweight and efficient, while still maintaining similar or even better performance on NLP tasks. It was achieved by applying two main techniques: factorization and cross-layer parameter sharing.

Factorization of the transformer-based architectures was used to reduce the number of parameters in the model, which makes the training and inference faster and more memory-efficient. Cross-layer parameter sharing was used to share the parameters across the layers, which also reduces the number of parameters. ALBERT is available in different versions such as ALBERT-base and ALBERT-xxlarge, which can be fine-tuned to specific tasks and datasets. It has been proven to have better performance on many NLP tasks compared to BERT and it is widely used in natural language understanding and text classification tasks.

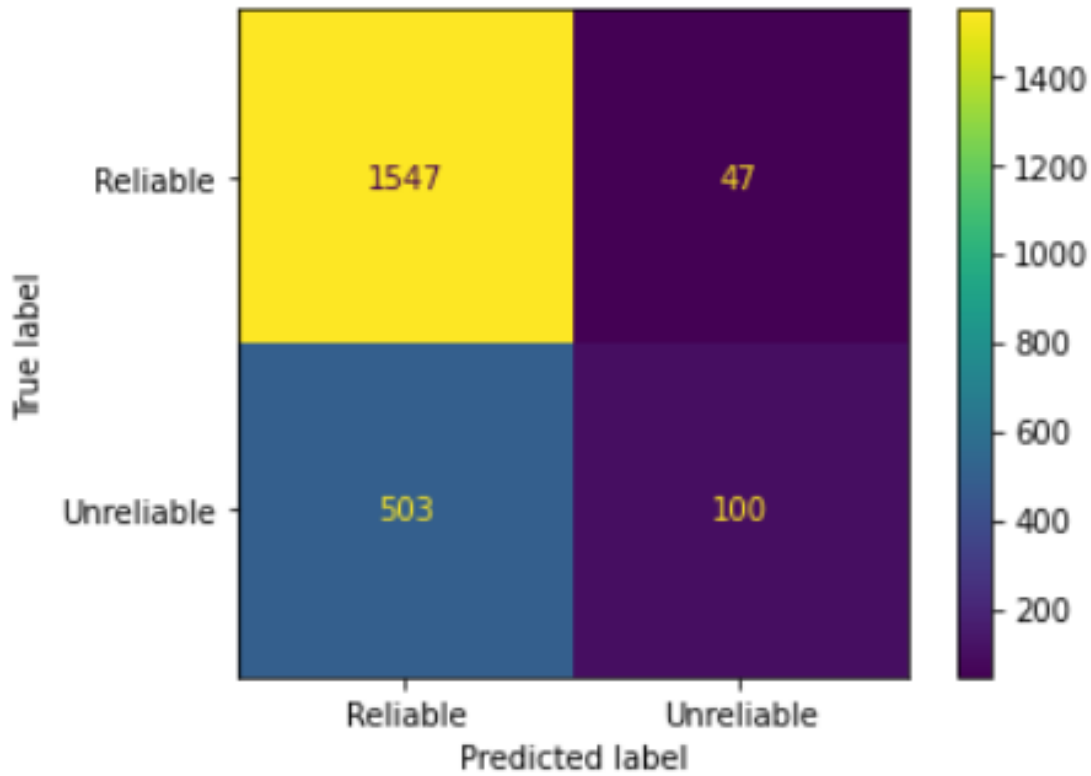
In summary, ALBERT is a light version of BERT, a pre-trained transformer-based neural network model for natural language processing tasks that is more lightweight and efficient while still maintaining similar or even better performance on NLP tasks. It was achieved by applying factorization and cross-layer parameter-sharing techniques. It is widely used in natural language understanding and text classification tasks.

ALBERT Model Summary:

Layer (type)	Output Shape	Param #	Connected to
text (InputLayer)	[(None,)]	0	[]
keras_layer (KerasLayer)	{'input_mask': (None, 128), 'input_word_ids': (None, 128), 'input_type_ids': (None, 128)}	0	['text[0][0]']
keras_layer_1 (KerasLayer)	{'pooled_output': (None, 768), 'default': (None, 768), 'sequence_output': (None, 128, 768), 'encoder_outputs': [(None, 128, 768), (None, 128, 768), (None, 128, 768), (None, 128, 768), (None, 128, 768), (None, 128, 768), (None, 128, 768), (None, 128, 768), (None, 128, 768), (None, 128, 768), (None, 128, 768), (None, 128, 768)]}	11683584	['keras_layer[0][0]', 'keras_layer[0][1]', 'keras_layer[0][2]']
output (Dense)	(None, 1)	769	['keras_layer_1[0][13]']
Total params: 11,684,353			
Trainable params: 769			
Non-trainable params: 11,683,584			



	precision	recall	f1-score	support
0	0.75	0.97	0.85	1594
1	0.68	0.17	0.27	603
accuracy			0.75	2197
macro avg	0.72	0.57	0.56	2197
weighted avg	0.73	0.75	0.69	2197



Comparison Table

Models	Accuracy (Training Dataset)	Accuracy (Testing Dataset)
LSTM	99.75%	81.30%
GRU	98.54%	78.98%
CNN	98.59%	81.36%
CNN + LSTM	98.59%	82.80%
BERT	73.67%	75%
ALBERT	72.33%	75%